

Denoising with a learned DNN prior

We consider the classical image or data denoising problem, where the goal is to remove zero-mean white Gaussian noise from a given image or data point. In more detail, our goal is to obtain an estimate of a vector $y_0 \in \mathbb{R}^n$ from the noisy observation

$$y = y_0 + \eta,$$

where η is zero-mean Gaussian noise with covariance matrix σ^2/nI , and y_0 lies in the range of the generator, i.e., $y_0 = G(x_0)$.

We consider the following two-step denoising algorithm:

1. Obtain an estimate \hat{x} of the latent representation by minimizing the empirical loss

$$f(x) = \|G(x) - y_0\|_2^2$$

using gradient descent.

2. Obtain an estimate of the image as $\hat{y} = G(\hat{x})$.

We learn G by training an encoder-decoder network (i.e., an autoencoder) and taking G as the decoder.

```
In [1]: import torch
import torch.nn as nn
import torch.utils as utils
from torch.autograd import Variable
import torchvision.datasets as dset
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
%matplotlib inline
import torch.nn.functional as F

import random
import numpy as np
import collections
```

Get data

```
In [2]: mnist_train = dset.MNIST("./", train=True, transform=transforms.ToTensor(), target_transform=None, download=True)
mnist_test = dset.MNIST("./", train=False, transform=transforms.ToTensor(), target_transform=None, download=True)

batch_size = 1
train_set = [ex for ex in torch.utils.data.DataLoader(dataset=mnist_train, batch_size=batch_size, shuffle=True)]
test_set = [ex for ex in torch.utils.data.DataLoader(dataset=mnist_test, batch_size=batch_size, shuffle=True)]

# construct training and test set only consisting of twos
def extract_nu(dset, nu):
    eset = []
    for image, label in train_set:
        if label.numpy() == nu:
            eset.append((image, label))
    return eset
train_twos = extract_nu(train_set, 2)
test_twos = extract_nu(test_set, 2)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./MNIST/raw/train-images-idx3-ubyte.gz
Failed to download (trying next):
HTTP Error 503: Service Unavailable

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz to ./MNIST/raw/train-images-idx3-ubyte.gz

Extracting ./MNIST/raw/train-images-idx3-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Failed to download (trying next):
HTTP Error 503: Service Unavailable

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz to ./MNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./MNIST/raw/train-labels-idx1-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Failed to download (trying next):
HTTP Error 503: Service Unavailable

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz to ./MNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./MNIST/raw/t10k-images-idx3-ubyte.gz to ./MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Failed to download (trying next):
HTTP Error 503: Service Unavailable

Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz
Downloading https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz to ./MNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./MNIST/raw/t10k-labels-idx1-ubyte.gz to ./MNIST/raw

Processing...
Done!

/Users/murong/.conda/envs/pythonProject/lib/python3.8/site-packages/torchvision/datasets/mnist.py:502: UserWarning: The given NumPy array is not writeable, and PyTorch does not support non-writeable tensors. This means you can write to the underlying (supposedly non-writeable) NumPy array using the tensor. You may want to copy the array to protect its data or make it writeable before converting it to a tensor. This type of warning will be suppressed for the rest of this program. (Triggered internally at /Users/dis-tiller/project/conda/conda-bld/pytorch_1616554845587/work/torch/csrc/autograd/utils/tensor_numpy.cpp:143.)
  return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

Specification of the autoencoder

Below, specify the encoder and decoder parameterized by lengthbottleneck:

- Encoder: Two-layer linear neural network consisting of linear layer + relu + linear layer + relu. The first linear layer takes the 784 input neurons, maps it to 400 neurons, and the second layer maps the 400 neurons to lengthbottleneck neurons.
- Decoder: linear layer + relu + linear layer + relu; lengthbottleneck neurons -> 400 neurons -> 784 neurons

```
In [3]: class Encoder(nn.Module):
        def __init__(self, lenbottleneck = 20):
            super(Encoder, self).__init__()
            self.fc1 = nn.Linear(28*28, 20*20)
            self.fc2 = nn.Linear(20*20, lenbottleneck)

        def forward(self, x):
            x = x.view(-1, 28*28)
            x = F.relu(self.fc1(x))
            x = F.relu(self.fc2(x))
            return x

class Decoder(nn.Module):
    def __init__(self, lenbottleneck = 20):
        super(Decoder, self).__init__()
        self.fc1 = nn.Linear(lenbottleneck, 20*20)
        self.fc2 = nn.Linear(20*20, 28*28)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = x.view(28, 28)
        return x

In [4]: def train(encoder, decoder, trainset, learning_rate = 0.001, epoch = 10):
        parameters = list(encoder.parameters()) + list(decoder.parameters())
        loss_func = nn.MSELoss() # mean square loss after the decoder
        print("learning_rate: ", learning_rate)
        optimizer = torch.optim.SGD(parameters, lr=learning_rate)
        ctr = 0
        for i in range(epoch):
            print("at epoch ", i + 1, "/", epoch)
            for (image, label) in trainset:
                ctr += 1
                image = Variable(image)
                optimizer.zero_grad()
                output = encoder(image)
                output = decoder(output)
                loss = loss_func(output, image)
                loss.backward()
                optimizer.step()
            print("trained in ", ctr, " iterations")
        return encoder, decoder
```

Training autoencoders

```
In [5]: # train autoencoder for all digits
K = 10
encoder10 = Encoder(K)
decoder10 = Decoder(K)
encoder, decoder = train(encoder10, decoder10, train_set[:10000], 1.0)

learning_rate: 1.0
at epoch 1 / 10

/Users/murong/.conda/envs/pythonProject/lib/python3.8/site-packages/torch/nn/modules/loss.py:528: UserWarning: Using a target size
(torch.Size([1, 1, 28, 28])) that is different to the input size (torch.Size([28, 28])). This will likely lead to incorrect result
s due to broadcasting. Please ensure they have the same size.
  return F.mse_loss(input, target, reduction=self.reduction)

at epoch 2 / 10
at epoch 3 / 10
at epoch 4 / 10
at epoch 5 / 10
at epoch 6 / 10
at epoch 7 / 10
at epoch 8 / 10
at epoch 9 / 10
at epoch 10 / 10
trained in 100000 iterations

In [6]: # train autoencoder for all digits
K = 20
encoder20 = Encoder(K)
decoder20 = Decoder(K)
encoder20, decoder20 = train(encoder20, decoder20, train_set[:10000], 1.0)


learning_rate: 1.0
at epoch 1 / 10
at epoch 2 / 10
at epoch 3 / 10
at epoch 4 / 10
at epoch 5 / 10
at epoch 6 / 10
at epoch 7 / 10
at epoch 8 / 10
at epoch 9 / 10
at epoch 10 / 10
trained in 100000 iterations
```

Check output of autoencoder: print a few input/output pairs

```
In [7]: # given two lists of images as np-arrays, plot them as a row
def plot_images(top,bottom):
    fig, axes = plt.subplots(nrows=2, ncols=10, sharex=True, sharey=True, figsize=(20,4))
    for images, row in zip([top, bottom], axes):
        for img, ax in zip(images, row):
            ax.imshow(img, cmap='Greys_r')
            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)
    fig.tight_layout(pad=0.1)
    return fig
```

```
In [8]: in_imgs = []
out_imgs = []
numprint = 10
for i,(img,label) in enumerate(test_set):
    if i >= numprint:
        break
    out_img = decoder20(encoder20(Variable(img)))
    in_imgs += [img[0][0].numpy()] # img is 1x1x28x28 tensor
    out_img.data.clamp_(0, 1)
    out_imgs += [out_img.data.numpy()]

fig = plot_images(in_imgs,out_imgs)
```



Task 1

Denoising with a trained decoder

Write the function for denoising below that denoises a noisy image given a generator learned earlier that maps a K-dimensional input space to an image.


```
In [9]: # denoise by recovering estimating a latent representation and passing that through the decoder
def denoise(net,noisy_image,K):
    # Step 1: Estimate the latent representation z
    z = torch.rand(K, requires_grad=True, device='cpu')
    learning_rate = 1.0
    loss_func = nn.MSELoss()
    optimizer = torch.optim.SGD([z], lr=learning_rate)
    net.training = False
    ite = 0
    diff = 1e4
    loss_prev = 1e4
    tol = 1e-6
    while diff > tol:
        ite += 1
        optimizer.zero_grad()
        decoded = net(z)
        loss = loss_func(decoded, noisy_image)
        loss.backward()
        optimizer.step()
        diff = abs(loss_prev - loss)
        loss_prev = loss
    print("trained in ", ite, " iterations to find the latent representation")
    # Step 2: Decode the latent estimation
    recovered_img = net(z)
    return recovered_img
```

Visualize denoising performance

```
In [12]: K = 20
Sigmas = [0.3*i for i in range(10)] # noise variances
noisy_imgs = []
rec_imgs = []
for (img,label),sigma in zip(test_set[:len(Sigmas)],Sigmas):
    img = img[0][0]
    noise = np.sqrt(sigma)*torch.norm(img)/np.sqrt(28*28)*torch.randn(28, 28)
    noisy_img = img + noise
    rec_img = denoise(decoder20,Variable(noisy_img),K)
    rec_img.data.clamp_(0, 1)
    noisy_imgs += [noisy_img.numpy()]
    rec_imgs += [rec_img.data.numpy()]

# plot and save to file
fig = plot_images(noisy_imgs,rec_imgs)
fig.savefig("denoising_ex_0.1.png")
```

```
trained in 1645 iterations to find the latent representation
trained in 2095 iterations to find the latent representation
trained in 1164 iterations to find the latent representation
trained in 2431 iterations to find the latent representation
trained in 2182 iterations to find the latent representation
trained in 1055 iterations to find the latent representation
trained in 1430 iterations to find the latent representation
trained in 708 iterations to find the latent representation
trained in 986 iterations to find the latent representation
trained in 2324 iterations to find the latent representation
```



Task 2

Compare empirically which network (the one with K=10, or the one with K=20). Your choice how to do this comparison.


```
In [13]: def check_denosig_performance(clean_imgs, rec_imgs):
    """
    Use average MSE of reconstructed images as the performance criterion.
    """
    mse = []
    for i in range(len(clean_imgs)):
        mse.append(np.square(np.subtract(clean_imgs[i], rec_imgs[i])).mean())
    return np.mean(mse)
```

```
In [18]: K_10 = 10
K_20 = 20
Sigmas = [0.3 * i for i in range(10)] # noise variances
noisy_imgs = []
clean_imgs = []
rec_imgs_10 = []
rec_imgs_20 = []
for (img, label), sigma in zip(test_set[:len(Sigmas)], Sigmas):
    img = img[0][0]
    noise = np.sqrt(sigma) * torch.norm(img) / np.sqrt(28 * 28) * torch.randn(28, 28)
    noisy_img = img + noise
    rec_img_10 = denoise(decoder10, Variable(noisy_img), K_10)
    rec_img_20 = denoise(decoder20, Variable(noisy_img), K_20)
    rec_img_10.data.clamp_(0, 1)
    rec_img_20.data.clamp_(0, 1)
    noisy_imgs += [noisy_img.numpy()]
    clean_imgs += [img.numpy()]
    rec_imgs_10 += [rec_img_10.data.numpy()]
    rec_imgs_20 += [rec_img_20.data.numpy()]
```

```
trained in 1227 iterations to find the latent representation
trained in 1530 iterations to find the latent representation
trained in 1620 iterations to find the latent representation
trained in 1924 iterations to find the latent representation
trained in 696 iterations to find the latent representation
trained in 1244 iterations to find the latent representation
trained in 1545 iterations to find the latent representation
trained in 2441 iterations to find the latent representation
trained in 1816 iterations to find the latent representation
trained in 1695 iterations to find the latent representation
trained in 1429 iterations to find the latent representation
trained in 872 iterations to find the latent representation
trained in 1065 iterations to find the latent representation
trained in 1025 iterations to find the latent representation
trained in 1948 iterations to find the latent representation
trained in 1273 iterations to find the latent representation
trained in 1139 iterations to find the latent representation
trained in 2014 iterations to find the latent representation
trained in 1709 iterations to find the latent representation
trained in 1290 iterations to find the latent representation
```


```
In [19]: # plot and save to file
print('Reconstruction performance of using network_10: ')
fig_10 = plot_images(noisy_imgs, rec_imgs_10)
fig_10.savefig("decoder10.png")
```

Reconstruction performance of using network_10:



```
In [20]: print('Reconstruction performance of using network_20: ')
fig_20 = plot_images(noisy_imgs, rec_imgs_20)
fig_20.savefig("decoder20.png")
```

Reconstruction performance of using network_20:



```
In [21]: mse_10 = check_denosing_performance(clean_imgs, rec_imgs_10)
mse_20 = check_denosing_performance(clean_imgs, rec_imgs_20)
print('The average MSE of network_10 is ', mse_10)
print('The average MSE of network_20 is ', mse_20)
```

The average MSE of network_10 is 0.015664935
The average MSE of network_20 is 0.012662632

Because the average reconstruction MSE of network_20 is smaller than network_10, using the latent dimension K=20 is better for this denoising task.